

Lab Handout 6: ThreadPool and Networking

Students are encouraged to share their ideas in the [#lab6](#) Slack channel. SCPD students are welcome to reach out to me or Hemanth directly if they have questions that can't be properly addressed without being physically present for a discussion section. The lab checkoff sheet for all students—both on-campus and off—can be found [right here](#).

Problem 1: ThreadPool Thought Questions

- Presented below are partial implementations of my own `ThreadPool::wait` and my `ThreadPool::worker` method.

```
void ThreadPool::worker(size_t workerID) {
    while (true) {
        // code here waits for a thunk to be supplied and then executes it
        lock_guard<mutex> lg(lock);
        if (--count == 0) done.notify_all();
    }
}

void ThreadPool::wait() {
    lock_guard<mutex> lg(lock);
    done.wait(lock, [this]{ return count == 0; });
}

void ThreadPool::worker(size_t workerID) {
    while (true) {
        // code here waits for a thunk to be supplied and then executes it
        lock_guard<mutex> lg(lock);
        if (--count == 0) done.notify_all();
    }
}
```

- Briefly describe a simple `ThreadPool` test program that would have deadlocked had `ThreadPool::worker` called `done.notify_one` instead of `done.notify_all`.

- Assume `ThreadPool::worker` gets through the call to `done.notify_all` and gets swapped off the processor immediately after the `lock_guard` is destroyed. Briefly describe a situation where a thread that called `ThreadPool::wait` still won't advance past the `done.wait` call.
- Had `done.notify_all` been called unconditionally (i.e. not just because `count` was zero), the `ThreadPool` would have still worked correctly. Why is this true, and why is the `if` test still the right thing to include?
- Consider the two server implementations below, where the sequential `handleRequest` function always takes exactly 1.500 seconds to execute. The two servers would respond very differently if 1000 clients were to connect—one per 1.000 seconds—over a 1000 second window. What would the 500th client experience when it tried to connect to the first server? What would the 500th client experience when it tried to connect to the second?

```
// Server Implementation 2
int main(int argc, char *argv[]) {
    ThreadPool pool(1);
    int server = createServerSocket(12346); // sets the backlog to 128
    while (true) {
        int client = accept(server, NULL, NULL);
        pool.schedule([client] { handleRequest(client); });
    }
}

// Server Implementation 1
int main(int argc, char *argv[]) {
    int server = createServerSocket(12345); // sets the backlog to 128
    while (true) {
        int client = accept(server, NULL, NULL);
        handleRequest(client);
    }
}

// Server Implementation 2
int main(int argc, char *argv[]) {
    ThreadPool pool(1);
    int server = createServerSocket(12346); // sets the backlog to 128
    while (true) {
```

```

    int client = accept(server, NULL, NULL);
    pool.schedule([client] { handleRequest(client); });
}
}

```

Problem 2: Networking, Client/Server, Request/Response

- Explain the differences between a pipe and a socket.
- Explain how system calls are a form of client/server and request/response.
- Describe how networking is just another form of function call and return. What “function” is being called? What are the parameters? And what's the return value? Assume HTTP is the operative protocol.
- Describe the network architecture needed for:
 - Email servers and clients
 - Peer-to-Peer Text Messaging via cell phone numbers
 - Skype
- As it turns out, each of the three network applications above all make use of custom protocols. Which ones could have relied on HTTP and/or HTTPS instead of custom protocols?

Problem 3: Hello Server Short Answers

Consider the following server called `hello-server`:

```

static void handleRequest(int client) {
    sockbuf sb(client);
    iosockstream ss(&sb);
    ss << "Hello from server (pid: " << getpid() << ")" << endl;
}

static void runServer(int server) {
    cout << "Firing up server inside process (pid: " << getpid() << ")." << endl;
    while (true) {
        int client = accept(server, NULL, NULL);
    }
}

```

```

    cout << "Request handled by server (pid " << getpid() << ")." << endl;
    handleRequest(client);
}
}

int main(int argc, char *argv[]) {
    int server = createServerSocket(33334);
    for (size_t i = 0; i < kNumWorkers; i++) {
        if ((workerPIDs[i] = fork()) == 0) runServer(server);
    }
    signal(SIGINT, shutdownAllServers);
    runServer(server); // let master process be consumed by same server

    return 0;
}static const size_t kNumWorkers = 3;
static pid_t workerPIDs[kNumWorkers];
static void shutdownAllServers(int unused) {
    for (size_t i = 0; i < kNumWorkers; i++) {
        cout << "Shutting down worker (pid " << workerPIDs[i] << ")." << endl;
        kill(workerPIDs[i], SIGINT);
        waitpid(workerPIDs[i], NULL, 0);
    }
    cout << "Shutting down orchestrator." << endl;
    exit(0);
}

static void handleRequest(int client) {
    sockbuf sb(client);
    iosockstream ss(&sb);
    ss << "Hello from server (pid: " << getpid() << ")" << endl;
}

static void runServer(int server) {
    cout << "Firing up server inside process (pid: " << getpid() << ")." << endl;

```

```

while (true) {
    int client = accept(server, NULL, NULL);
    cout << "Request handled by server (pid " << getpid() << ")." << endl;
    handleRequest(client);
}
}

int main(int argc, char *argv[]) {
    int server = createServerSocket(33334);
    for (size_t i = 0; i < kNumWorkers; i++) {
        if ((workerPIDs[i] = fork()) == 0) runServer(server);
    }
    signal(SIGINT, shutdownAllServers);
    runServer(server); // let master process be consumed by same server

    return 0;
}

```

If I launch the above executable on myth4, hit it 12 times by repeatedly typing `telnet myth4 33334` at the prompt from a myth22 shell, and then press `ctrl-c` on myth4, I get the following:

```

myth4> ./hello-server
Firing up hello server inside process with pid 19941.
Firing up hello server inside process with pid 19940.
Firing up hello server inside process with pid 19942.
Firing up hello server inside process with pid 19943.
Request handled by server (pid 19940).
Request handled by server (pid 19942).
Request handled by server (pid 19941).
Request handled by server (pid 19943).
Request handled by server (pid 19940).
Request handled by server (pid 19942).
Request handled by server (pid 19941).
Request handled by server (pid 19943).

```

```
Request handled by server (pid 19940).
Request handled by server (pid 19942).
Request handled by server (pid 19941).
Request handled by server (pid 19943).
^CShutting down server with pid 19941.
Shutting down server with pid 19942.
Shutting down server with pid 19943.
Shutting down master.
```

Based on the code and test run you see above, answer each of the following five short answer questions.

- Note the very first line of the server's `main` function creates a server socket that listens to 33334 for incoming network activity. Why, after the `for` loop within `main`, are all of the child processes also listening to port 33334 through the same server socket descriptor?
- During lecture, we referred to port numbers as virtual process ids. What did we mean by that, and why are virtual process ids needed with networked applications?
- What happens if the call to `exit(0)` is removed from the implementation of `shutdownServers` and we send a `SIGINT` to the master process in the suite of 4 `hello-server` servers?
- If the call to the `signal` function is moved to reside above the `for` loop in `main` instead of below it, does that impact our ability to close down the full suite of 4 `hello-server` servers?
- The above program doesn't close the server sockets in `shutdownServers`. Describe a simple way to ensure a server socket is properly closed before the surrounding server executable exits.

Problem 4: Web Servers and Python Scripts

It's high time you implement your own multithreaded HTTP web server, making use of a `ThreadPool` to manage the concurrency, and all of the multiprocessing functions (`fork` and all related functions) to launch another executable—in all cases, a Python script—to ingest the majority of the HTTP request from the client and publish the entire HTTP response back to it.

Here are the pertinent details:

- The web server listens to port 80 like all normal web servers do.
- Every single HTTP request is really a request to invoke a python script. For instance, if the first line of a request is `GET /add.py?one=11&two=50 HTTP/1.0`, the web server exe-

cutes the {"python", "./scripts/add.py", "GET", "/add.py?one=11&two=50", "HTTP/1.0"} argument vector.

- All python scripts are housed within the `scripts` directory. `GET /one.py` invokes `./scripts/one.py`, `GET /deep/down/below.py` invokes `./scripts/deep/down/below.py`, and so forth.
- All python scripts expect to read the entire HTTP request (save for the first line) in through standard in, and the entire HTTP response is published via standard out.
- All python scripts accept three arguments, which are the tokens that make up the first line of the HTTP request. The web server itself must read in the first line from the client, tokenize it so it knows what to pass as arguments to the relevant python script, and then assume the python script pulls everything else.
- If the URL identifies a resource whose name doesn't end up ".py", then your server can simply ignore the request and close the connection.
- Similarly, if the URL identifies a resource whose name ends in ".py", but the python script doesn't exist, your web server can ignore the request and close the connection.
- If the URL identifies a resource whose name ends in ".py" but includes "..", your web server can ignore the request and close the connection. (This is presumably an attempt to reach up above the `scripts` directory and gain access to an arbitrary executable.)
- You may not use the `system` function as you're permitted in your map-reduce assignment (which goes out tonight). You must use `fork`, `execvp`, and other related functions.
- You may assume that all python scripts successfully invoked always succeed, so you needn't check any status codes via `WIFEXITED`, `WEXITSTATUS`, etc.
- Your implementation must close all unused descriptors and cull all zombie processes.
- Your implementation should not make use of any signal handlers.

Your implementation can make use of the following routine, which reads the first line of an HTTP request coming over the provided client socket (up to and including the "\r\n"), and surfaces the method (e.g. "GET"), the full URL ("/add.py?one=11&two=50"), the path ("/add.py"), and the protocol ("HTTP/1.1") through the four `strings` shared by reference.

```
static void parseRequest(int client, string& method, string& url, string& path,
string& protocol);
```

You're to complete the implementation of the web server by fleshing out the details of the `handleRequest` function.

```
static void handleRequest(int client) {
    int main() {
        int server = createServerSocket(/* port = */ 80);
```

```
    runServer(server);  
    return 0; // never gets here, but compiler doesn't know that  
}  
  
static const kNumThreads = 16;  
static void runServer(int server) {  
    ThreadPool pool(kNumThreads);  
    while (true) {  
        int client = accept(server, NULL, NULL);  
        pool.schedule([client] { handleRequest(client); });  
    }  
}  
  
static void handleRequest(int client) {
```